

Top 5 supercomputers

1. Frontier

Size: Frontier is expected to have over 100,000 nodes, each equipped with multiple processors.

Location: Oak Ridge National Laboratory (ORNL) in Tennessee, USA.

What it's for: Frontier is designed for various scientific applications such as climate modeling, materials science, and energy research.

Who's working on it: Frontier is developed by Cray Inc., a part of Hewlett Packard Enterprise.

Cost: Estimated to be around \$600 million.

Status: Became operational in 2022. Peak Performance 1.5 exaflops.

2. Fugaku

Size: Fugaku has over 7.3 million processor cores.

Location: RIKEN Center for Computational Science in Kobe, Japan.

What it's for: Fugaku is used for applications like drug discovery, weather forecasting, and scientific simulations.

Who's working on it: Jointly developed by RIKEN and Fujitsu.

Cost: Estimated to be around \$1 billion.

Status: Became operational in 2021. Peak Performance 442 petaflops.

3. Lumi

Size: Lumi is expected to have tens of thousands of nodes with multiple processors.

Location: CSC - IT Center for Science in Kajaani, Finland.

What it's for: Used for addressing scientific and industrial challenges such as climate modeling and drug discovery.

Who's working on it: Developed by a consortium led by the Finnish IT Center for Science.

Cost: Estimated to be around €200 million.

Status: Became operational in 2022. Peak Performance 309 petaflops.

4. Leonardo

Size: Leonardo is expected to have several thousand nodes with multiple processors.

Location: Italian supercomputing center CINECA in Bologna, Italy.

What it's for: Supporting research in fields like astrophysics and computational biology.

Who's working on it: Developed by CINECA.

Cost: Estimated to be around €120 million.

Status: Became operational in 2023. Peak Performance 174 petaflops.

5. Summit

Size: Summit has over 2.4 million processor cores.

Location: Oak Ridge National Laboratory (ORNL) in Tennessee, USA.

What it's for: Used for scientific research including climate modeling and materials science.

Who's working on it: Developed by IBM, in partnership with NVIDIA, Mellanox, and others.

Cost: Estimated to be around \$200 million.

Status: Became operational in 2018. Peak Performance 148.8 petaflops